Evaluation results

The script evaluate.py runs a fixed set of requests and checks which tools the LLM called. Decoding is greedy, so the same harness gives the same score on every run, and a change in score comes from a change in the harness.

Change	Requests passed
First system prompt	1 of 5
Rule: one tool call per thing asked for	3 of 5
Hint about requests that need several tools (a sixth request added)	3 of 6
Score only the tool calls that actually ran	5 of 6
Router scores each clause of the request	6 of 6

Each edit to the system prompt fixed one request and broke another. The final fix was in the router, not in the prompt. After the completion check was added for the 40-second narrated video, the score stayed at 6 of 6.